

Reinforcement Learning With Reward Machines in Stochastic Games

Reinforcement Learning Course - Final Project — 06-2026

Matteo Liotta

`matteo.liotta@studenti.units.it`

Mattia Simonutti

`mattia.simonutti@studenti.units.it`

University of Trieste

Department of Mathematics and Geosciences

Problem Setting

- **Goal:** Learn to the optimal best-response strategy at a **Nash equilibrium** for each agent in a **general-sum game setting**.
- **Problem:** If multi-agent complex stochastic games induce system **non-Markovianity**, how to handle **non-Markovian** reward functions?
- **Proposed Approach:** Usage of **Reward Machines** to incorporate the high-level game knowledge and **converting non-Markovian reward functions** into **standard Markovian ones** [1].

Basic Definitions

Stochastic Games with Non-Markovian Rewards

Definition — Two-Player General-Sum Stochastic Game (SG)

Modeled as the tuple $\mathcal{G} = (S, s_I, A_e, A_a, p, R_e, R_a, \gamma)$, where:

- **Finite State Space (S):** consisting of **states** $s = [s_e; s_a]$, with s_e, s_a substates for *Ego* and *Adv* agents respectively. $s_I \in S$ is the **initial** state.
- **Action Spaces:** Each agent has its own finite set of actions A_e and A_a .
- **Probabilistic Transition Function:** $p : S \times A_e \times A_a \times S \rightarrow [0, 1]$, i.e. $p(s' | s, a_e, a_a)$
- **Non*-Markovian Reward functions:** $R_e \text{ or } a : (S \times A_e \times A_a)^+ \times S \rightarrow \mathbb{R}$ map from **complete trajectory histories** to a real value.

(*) Usual definition assumes system Markovianity.

Trajectories and Rewards

- **Trajectory** (at step k): sequence of **states** and **actions** defined as

$$t_{0:k} = s_0 a_{\epsilon,0} a_{\alpha,0} s_1 \dots s_{k-1} a_{\epsilon,k-1} a_{\alpha,k-1} s_k \quad \text{with} \quad s_0 = s_I$$

- **Reward Sequence**: Each agent $i \in \{\epsilon, \alpha\}$ receives a sequence $r_{i,1} \dots r_{i,k}$, where each value is determined by the **history up to that step**:

$$r_{i,j} = R_i(s_0 a_{\epsilon,0} a_{\alpha,0} \dots s_j a_{\epsilon,j} a_{\alpha,j} s_{j+1}) = R_i(t_{0:j+1}) \quad \text{for each} \quad j \leq k$$

- **Shared Observation**: Both agents **observe the same** trajectory.
- **Discounted Cumulative Reward**: Given a trajectory $t_{0:k}$, agents achieves $\sum_{j=1}^k \gamma^{j-1} r_{i,j}$, where $\gamma \in (0, 1]$

Strategies and Objectives

- **Markovian Strategy:** Agents are equipped with a $\pi_i : S \times A_i \rightarrow [0, 1]$ $i \in \{\epsilon, \alpha\}$ strategy, with $\pi_i(a_i|s)$ probability of agent i selecting action a_i from state s .
- **Objective:** Maximize the **expected** discounted cumulative reward given policy of both agents, from both point of view.

Definition — SG Expected Discounted Cumulative Reward

Given agent $i \in \{\epsilon, \alpha\}$, the expected DCR under policies π_ϵ and π_α is defined as:

$$\tilde{v}_i(s, \pi_\epsilon, \pi_\alpha) = \sum_{t=0}^{\infty} \gamma^t \mathbb{E}(r_{i,t} | \pi_\epsilon, \pi_\alpha, s_0 = s)$$

Nash Equilibrium

In Stochastic Games, each agent's strategy is a **best-response to the other agent's strategy**, if each **cannot improve its own reward** by changing its strategy, given a **fixed opponent's policy**.

Definition — Nash Equilibrium in Stochastic Games

For a stochastic game $\mathcal{G} = (S, s_I, A_e, A_a, R_e, R_a, \gamma, p)$, the strategies of the ego agent, π_e^* , and the adversarial agent, π_a^* , are at a **Nash equilibrium** if for any π_e, π_a :

$$\begin{aligned}\tilde{v}_e(s, \pi_e^*, \pi_a^*) &\geq \tilde{v}_e(s, \pi_e, \pi_a^*) \\ \tilde{v}_a(s, \pi_e^*, \pi_a^*) &\geq \tilde{v}_a(s, \pi_e^*, \pi_a)\end{aligned}$$

With $*$ defining a fixed point of stability

Definition — Reward Machine (RM)

Mealy machine ($\approx$ DFA) designed to **encode non-Markovian rewards** with **state transitions** and **output functions**.

$$\mathcal{A} = (V, v_I, 2^{\mathcal{P}}, \mathbb{R}, \delta, \sigma)$$

- **Set of RM states (V):** Finite collection of internal **game** evolution tracking **states**, with v_I as the **initial** state.
- **Input Alphabet ($2^{\mathcal{P}}$):** Set of possible combinations of **propositional variables** of $p \in \mathcal{P}$, which encodes high-level environmental knowledge/events.
- **Transition function: Deterministic** game state progression $\delta : V \times 2^{\mathcal{P}} \rightarrow V$.
- **Output function:** $\sigma : V \times 2^{\mathcal{P}} \rightarrow \mathbb{R}$, associates values (**rewards**) along active trajectories.

Connecting SG and RMs

A link between RM and SG is required

Labelling Function (external to RMs)

$$L : S \times A_e \times A_a \times S \rightarrow 2^{\mathcal{P}}$$

$l_t = L(s_t, a_{a,t}, a_{e,t}, s_{t+1})$ is the **set of proposition** in $\mathcal{P}$ having **true** evaluation at step t .

RM can encode non-Markovian RF of a stochastic game

Given a SG, the RM $\mathcal{A}_i$ **encodes** the non-markovian reward function R_i of the game, meaning that **for each trajectory** $s_0 a_{e,0} a_{a,0} s_1 \dots s_k a_{e,k} a_{a,k} s_{k+1}$ (i.e. $l_0 l_1 \dots, l_k$) it outputs the **reward sequence** $r_{i,0} r_{i,1} \dots, r_{i,k}$.

I.e. the RM maps **any sequence of events to a sequence of rewards**.

Visual Examples of RMs and L

Reference paper focuses on **two-agent GS-SG** where each agent potentially given a different task to complete. where task completion **depends on the other's agent behavior**.




Figure. Example Grid World




Figure. Example Associated RM

$$L(s_t, a_a, a_e, s_{t+1}) = \{\text{red triangle}, \text{blue circle}\}$$

Equation. Example Labelling function output

Q-learning with RM for SG

Markovian formulation of SG with RMs

Definition — Stochastic Game with Reward Machines (SGRM)

Given an SG $\mathcal{G}$ and RMs $\mathcal{A}_e, \mathcal{A}_a$, a SGRM is defined as a product stochastic game:

$$\mathcal{H} = (S', s'_I, A_e, A_a, p', R'_e, R'_a, \gamma)$$

- **Augmented State Space:** $S' = S \times V_e \times V_a$. An **augmented state** is $s' = (s, v_e, v_a)$.
- **Transition Probabilities:** $p'(s', a_e, a_a, \bar{s}') = p(s, a_e, a_a, \bar{s})$ if $v'_i = \delta_i(v_i, L(s, a_e, a_a, \bar{s}))$ for both $i \in \{e, a\}$, and 0 otherwise.
- **Markovian Reward Functions:** $R'_i(s', a_e, a_a, \bar{s}') = \sigma_i(v_i, L(s, a_e, a_a, \bar{s}))$, $\forall i$

Why the SGRM formulation?

The newly introduced Stochastic Game $\mathcal{H}$ is proved to have a **Markovian rewards functions** R'_i **which expresses the original** Stochastic Game $\mathcal{G}$ **non-Markovian reward functions** R_i ($\forall i \in \{\mathbf{e}, \mathbf{a}\}$)

SGRM steps in practice

1. The agent consider S' to select actions:
 - 1.1 At each time step each agent **selects the action to perform in s , simultaneously** executed.
 - 1.2 Action execution causes **game state transition** from s to s' .

$$s \xrightarrow{a_e, a_a} s'$$

2. The tuple (s_t, a_e, a_a, s_{t+1}) becomes then available.
 - 2.1 With this input, L returns the **set of satisfied atomic proposition** of $\mathcal{P}$, l_t .
 - 2.2 Accordingly to available action with l_t , a **transition** is performed over each agent RM:

$$v_e \xrightarrow{l_t} v'_e \quad v_a \xrightarrow{l_t} v'_a$$

3. RM transitions lead to each agent reward obtainance: $r_{i,t} = \sigma(v_i, L(s, a_e, a_a, s'))$

Q-functions at Nash Equilibrium

In SGRM we define the **optima Q-function** for agent i the...

Q-function at a Nash Equilibrium

It is the **EDCR** obtained by agent i when both agents **follow a joint Nash equilibrium strategy** from the next period on. In augmented space $S \times V_e \times V_a$ it is defined:

$$q_i^*(s, v_e, v_a, a_e, a_a) = r_i + \gamma \sum_{s', v'} p(s', v' | s, v, a_e, a_a) \underbrace{\tilde{v}_i(s', v'_e, v'_a, \pi_e^*, \pi_a^*)}_{\text{Following NE from next period on}}$$

Finding optimal Policies

Given the both agent Q-function, **best response policies** are then originally extracted using the **Lemke-Howson method**:

$$\pi_{\epsilon}^{NE}, \pi_a^{NE} \leftarrow \text{LH_METHOD}(Q_{\epsilon}, Q_a)$$

...But an agent is blind to the other Q-function due to adversarial problem nature.

However, since both receive same signal $(s_t, a_{\epsilon}, a_a, s_{t+1})$, using **personal RMs**, they estimate opponents rewards, **estimating opponent Q-Function** too:

$$\underbrace{Q_{\epsilon, \epsilon} \quad Q_{a, \epsilon}}_{\text{Agent Ego estimations}} \quad \underbrace{Q_{a, a} \quad Q_{\epsilon, a}}_{\text{Agent Adv estimations}}$$

Finding optimal Policies (cont'd)

Definition: Stage Game

A two-player stage game is defined as (r_e, r_a) , where r_i is agent i 's reward function R_i , over the space of joint actions:

$$r_i = \{R_i(a_a, a_e) | a_a \in A_a, a_e \in A_e\}$$

In QRM-SG, a stage game is formulated at each time step based on current Q-functions in augmented state space.

Because the product stochastic game with reward machines has Markovian rewards, QRM-SG can apply Nash-Q updates: at each augmented state, it derives a Nash equilibrium of the current Q-value stage game and uses the associated best-response strategies in the Q-function update.

QSG-RM Algorithm Overview

QRM-SG Algorithm Key aspects

Actor Mechanics (Action Selection)

- ϵ -greedy policy is adopted to balance the **exploration** and **exploitation**.

$$a_i \leftarrow \begin{cases} \operatorname{argmax}_{a \in A_i} \pi_{ji}(a \mid s, v_e, v_a) & \text{with probability } 1 - \epsilon \\ \text{random} & \text{with probability } \epsilon \end{cases}$$

Critic Updates (Learning Loop)

Values are updated according to standard stochastic approximation:

$$q_{ji} \leftarrow (1 - \alpha_t) q_{ji} + \alpha_t (r_{i,t} + \gamma \bar{q}_{ji}(s', v'_e, v'_a))$$

where

$$\bar{q}_{ji} \leftarrow \pi_{e,i}(s', v'_e, v'_a) \pi_{a,i}(s', v'_e, v'_a) q_{ji}(s', v'_e, v'_a) \in \mathbb{R}$$

QRM-SG Phase 1 — Initialization

Process Breakdown (Lines 1–4)

Hyperparameter: episode length $eplength$, γ , ϵ

Input: Reward machines $\mathcal{A}_\epsilon, \mathcal{A}_\alpha$

$s \leftarrow InitialState()$; $v_\epsilon \leftarrow v_{I,\epsilon}$; $v_\alpha \leftarrow v_{I,\alpha}$

$q_{\epsilon i}(\cdot), q_{\alpha i}(\cdot) \leftarrow InitialQfunctions()$

Analysis:

- Each agent initializes **two** Q-tables to store values for self and opponent, allowing local equilibrium computation.
- SGRM (augmented) states are defined in the product space $S \times V_\epsilon \times V_\alpha$.

QRM-SG Phase 2 — Nash-Action Selection

Process Breakdown (Lines 7–13)

```
# for each episode, for each time step  $t$ :  
 $\bar{\epsilon} \leftarrow \text{GenerateRandomValue}(0, 1)$   
if  $\bar{\epsilon} < \epsilon$  then  
     $a_i \leftarrow \text{ChooseActionRandomly}()$  # Acting at random  
else  
    for  $i \in \{\epsilon, \alpha\}$  do  
         $\pi_{\epsilon i}, \pi_{\alpha i} \leftarrow \text{CalculateNashEquilibrium}(q_{\epsilon i}, q_{\alpha i})$  # Get current best strategy  
         $a_i \leftarrow \arg\max \pi_{ij}(a|s, v_{\epsilon}, v_{\alpha})$  # ... and select action accordingly  
    end for  
end if
```

Each agent solves a bimatrix stage game using the **LHmethod** to find mixed strategy equilibria.

Process Breakdown (Lines 14–18)

```
 $s' \leftarrow \text{ExecuteAction}(s, a_e, a_a)$  # Agent State transition  
 $l_t \leftarrow L(s, a_e, a_a, s')$   
for  $i \in \{e, a\}$  do  
     $v'_i \leftarrow \delta_i(v_i, l_t); r_{i,t} \leftarrow \sigma_i(v_i, l_t)$  # Game state transition and immediate rewards  
end for
```

- L bridges physical environmental transitions ($s \rightarrow s'$) with high level verification (l_t), allowing logic task transitions ($v_i \rightarrow v'_i$)
- Rewards $r_{i,t}$ are generated by the RM transitions, making them implicitly **history aware**.

QRM-SG Phase 4 - Nash Q-Learning

In single agent settings, Q-functions can be updated via the **off policy Q-learning**, with rule:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

But not if rewards depend on the joint behavior of more agents. The EDCR averaging w.r.t. policy π moves to:

$$\max_{a'} Q(s', a') \longrightarrow \mathbb{E}_{\pi_e, i, \pi_a, i} [Q_{e,i}(s', a')]$$

Nash-Q-Learning makes use of co-dependent expectation. Each agent i updates its estimates:

$$Q_{e,i}(s, \mathbf{a}) \leftarrow Q_{e,i}(s, \mathbf{a}) + \alpha(r_{e,i} + \gamma \mathbb{E}_{\pi_e, i, \pi_a, i} [Q_{e,i}(s', \mathbf{a}')] - Q_{e,i}(s, \mathbf{a}))$$

$$Q_{a,i}(s, \mathbf{a}) \leftarrow Q_{a,i}(s, \mathbf{a}) + \alpha(r_{a,i} + \gamma \mathbb{E}_{\pi_e, i, \pi_a, i} [Q_{a,i}(s', \mathbf{a}')] - Q_{a,i}(s, \mathbf{a}))$$

Current policies estimations are obtained with a call to **CalculateNashEquilibrium** function.

Compared to classical Q-Learning, the algorithm still looks ahead to the **next state** (s', v'_e, v'_a), but instead of looking at the entry with maximal value wrt action, it **averages with respect to all combinations of agents actions**, weighted by the **joint probability** of action choice.

Additional Keypoints

Introduction of the counterfactual Reasoning

Original paper Nash Q-learning is as a function of the augmented state space $S \times V_{\epsilon} \times V_{\alpha}$. Additionally, given a state, step update affects only the state $(s, v_{\epsilon}, v_{\alpha})$ after action a_{ϵ}, a_{α} .

Pointed out problem

Given a physical state s , we can reach next s' having **many different** RM internal states v_{ϵ}, v_{α} .

Since, the 4 Q-function updates highly depends on $|V_{\epsilon} \times V_{\alpha}|$, to perform a **general state game learning**, we would need to **experience a huge amount of steps/episodes**.

Introduction of the counterfactual Reasoning

However, physical states (s) and logical game states (v_e, v_a) are **independent**:

- Agents may move in gridworld and receive I_t **regardless of task related history**.
- RM states are meant to keep **only task relevant high-level informations** through v_i , **without spatial** nor **temporal** informations.

So, during each $s \rightarrow s'$ transition, we can **simulate** having **different RM states**, i.e. simulating **alternative historical backgrounds** at each $s \rightarrow s'$.

Counterfactual Reasoning

Instead of needing performing $s \rightarrow s'$ with every possible **history** permutation, we use experience all histories at once.

This shifts (*at its best*) complexity **away from exploration** and transfers it into the **computation of an individual trace**.

What does it means in practice?

For a $N \times M$ grid, 2 agents, 4 actions for agent, RMs with K_e , K_a states, if standard RM based code requires E episodes to converge, then with the counterfactual reasoning we can expect to converge in $\frac{E}{K_e \cdot K_a}$ episodes.

Counterfactual Reasoning (cont'd)

Counterfactual Reasoning

Since the **practical environment transition is independent from the RM state**, instead of updating the Q-value for state (s, v_e, v_a) **only**, we can

- Consider all RMs states $\tilde{v}_e, \tilde{v}_a$, simulate being in augmented state $(s, \tilde{v}_e, \tilde{v}_a)$.
- Simulate state and the RM transition to $\tilde{v}'_e, \tilde{v}'_a$ if s' is reached, i.e. reaching $(s', \tilde{v}'_e, \tilde{v}'_a)$
- Update the Q-function associated to $(s, \tilde{v}_e, \tilde{v}_a)$

Convergence Assumptions

Assumption 1: Exploration

Every state $s \in S$, RM states $v_e \in V_e$, $v_a \in V_a$, and actions $a_e \in A_e$, $a_a \in A_a$, are visited infinitely often when the number of episodes goes to infinity.

Assumption 2: Learning Rate

The learning rate α_t satisfies the following conditions for all t :

1. $0 < \alpha_t < 1$, $\sum_t \alpha_t(s, v_e, v_a, a_e, a_a) = \infty$, $\sum_t [\alpha_t(s, v_e, v_a, a_e, a_a)]^2 < \infty$, and the latter two hold uniformly and with probability 1.
2. $\alpha_t(s, v_e, v_a, a_e, a_a) = 0$ if $(s, v_e, v_a, a_e, a_a) \neq (s^t, v_e^t, v_a^t, a_e^t, a_a^t)$.

Assumption 3: Stage Game Structure

One of the following conditions holds during learning:

1. Every stage game $(q_{ei}(s, v_e, v_a), q_{ai}(s, v_e, v_a))$, $i \in \{e, a\}$, for all t, s, v_e, v_a , has a **global optimal point** $(\tilde{\pi}_{ei}, \tilde{\pi}_{ai})$ (i.e., $\tilde{\pi}_{ji}q_{ji} \geq \tilde{\pi}'_{ji}q_{ji}$ for any $\tilde{\pi}'_{ji}$, $j \in \{e, a\}$) and agents' rewards in this equilibrium are used to update their Q-functions.
2. Every stage game $(q_{ei}(s, v_e, v_a), q_{ai}(s, v_e, v_a))$, $i \in \{e, a\}$, for all t, s, v_e, v_a , has a **saddle point** $(\tilde{\pi}_{ei}, \tilde{\pi}_{ai})$ (i.e., $\tilde{\pi}_{ji}\tilde{\pi}_{-ji}q_{ji} \geq \tilde{\pi}'_{ji}\tilde{\pi}_{-ji}q_{ji}$ for any $\tilde{\pi}'_{ji}$, $\tilde{\pi}_{ji}\tilde{\pi}_{-ji}q_{ji} \leq \tilde{\pi}_{ji}\tilde{\pi}'_{-ji}q_{ji}$ for any $\tilde{\pi}'_{-ji}$, $j \in \{e, a\}$), and agents' rewards in this equilibrium are used to update their Q-functions.

Game Results and Experiments

Testing Environment: Evaluated on 6×6 (and 12×12) competitive, sparse-reward grid world variations of Pac-Man.

Comparative Baselines:

- Nash-Q: Pure spatial state-action tracking.

Case Study Empirical Analysis

- **Case Study I (Sequential Ordering):** QRM-SG converges to equilibrium by $\approx 7,500$ episodes, whereas unaugmented baselines fail entirely.
- **Case Study II & III (Energy Loops):** Explores complex recharge loops with randomized origin points. QRM-SG reaches stability within $\approx 1,000$ to $1,500$ runs.
- **Scalability:** Successfully finds the Nash equilibrium on an expanded 12×12 map space by episode 21,000.

Conclusion

Summary & Future Outlook

- **Core Contribution:** Bridges reward machines with MARL to model non-Markovian constraints in stochastic games.
- **Robust Architecture:** Augmented state transformations allow tabular actors to navigate complex environments successfully.
- **Future Directions:** Dynamically learning unknown reward machine configurations simultaneously during policy calculation; extensions to general-sum games with more agents.

References

- [1] Jueming Hu, Jean-Raphaël Gaglione, Yanze Wang, Zhe Xu, Ufuk Topcu, and Yongming Liu.

Reinforcement learning with reward machines in stochastic games.

arXiv preprint arXiv:2305.17372v3, cs.MA, 2023.