

Reinforcement Learning With Reward Machines in Stochastic Games

Reinforcement Learning Course - Final Project — 06-2026

Matteo Liotta

`matteo.liotta@studenti.units.it`

Mattia Simonutti

`mattia.simonutti@studenti.units.it`

University of Trieste

Department of Mathematics and Geosciences

Problem Setting

- **Problem:** How to tackle complex tasks in **multi-agent RL** where reward functions are **non-Markovian**?
- **Proposed Approach:** Usage of **Reward Machines** to incorporate the high-level knowledge of the task. By RM state incorporation, form an augmented product stochastic game, **converting non-Markovian rewards** into **standard Markovian ones** [1].
- **Goal:** Learn and converge to the optimal best-response strategy at a **Nash equilibrium** for each agent in a **general-sum game setting**.

Basic Definitions

Stochastic Games with Non-Markovian Rewards

Definition — Two-Player General-Sum Stochastic Game (SG)

Modeled as the tuple $\mathcal{G} = (S, s_I, A_e, A_a, p, R_e, R_a, \gamma)$, where:

- **Finite State Space (S):** consisting of states $s = [s_e; s_a]$, with s_e, s_a substates for *Ego* and *Adv* agents respectively. $s_I \in S$ is the initial state.
- **Action Spaces:** Each agent has its own finite set of actions A_e and A_a .
- **Probabilistic Transition Function:** $p : S \times A_e \times A_a \times S \rightarrow [0, 1]$
- **Non*-Markovian Reward functions:** $R_e \text{ or } a : (S \times A_e \times A_a)^+ \times S \rightarrow \mathbb{R}$ map from **complete trajectory histories** to real value domains, specifying the n-m rewards for both agents.

(*) Usual definition assumes system Markovianity.

Trajectories and Rewards

- **Trajectory:** Sequence of states and actions $t_{0:k+1} = s_0 a_{\epsilon,0} a_{\alpha,0} s_1 \dots s_k a_{\epsilon,k} a_{\alpha,k} s_{k+1}$ with $s_0 = s_I$.
- **Reward Sequence:** Given a **trajectory**, each agent ($i \in \{\epsilon, \alpha\}$) receives reward sequence $r_{i,1} \dots r_{i,k}$, where $r_{i,k} = R(s_0, a_{\epsilon,0}, a_{\alpha,0}, \dots, s_k)$.
- **Shared Observation:** Both agents **observe the same** trajectory.
- **Discounted Cumulative Reward:** Given a **trajectory** $t_{0:k+1}$, agents achieves a DCR reward $\sum_{t=0}^k \gamma^t r_{i,t}$
- **Finite episodes assumption:** The game is episodic with a finite horizon, ensuring that the cumulative reward is well-defined.

Strategies and Objectives

- **Markovian Strategy:** Agents are equipped with a $\pi_i : S \times A_i \rightarrow [0, 1]$ $i \in \{\epsilon, \alpha\}$ strategy, defining action given state selection probability.
- **Objective:** Maximize the **expected** discounted cumulative reward given policy of both agents, from both point of view.

Definition — SG Expected Discounted Cumulative Reward

Given agent $i \in \{\epsilon, \alpha\}$, the expected DCR under policies π_ϵ and π_α is defined as:

$$\tilde{v}_i(s, \pi_\epsilon, \pi_\alpha) = \sum_{t=0}^{\infty} \gamma^t \mathbb{E}(r_{i,t} | \pi_\epsilon, \pi_\alpha, s_0 = s)$$

Nash Equilibrium

Each agent's strategy is a **best-response to the other agent's strategy**, if each agent **cannot improve its own reward** by changing its strategy, given a **fixed opponent's policy**.

Definition — Nash Equilibrium in Stochastic Games

For a stochastic game $\mathcal{G} = (S, s_I, A_e, A_a, R_e, R_a, \gamma, p)$, the strategies of the ego agent, π_e^* , and the adversarial agent, π_a^* , are at a **Nash equilibrium** if:

$$\tilde{v}_e(s, \pi_e^*, \pi_a^*) \geq \tilde{v}_e(s, \pi_e, \pi_a^*)$$

$$\tilde{v}_a(s, \pi_e^*, \pi_a^*) \geq \tilde{v}_a(s, \pi_e^*, \pi_a)$$

hold for any π_e and π_a .

Objectives (cont'd)

The **actual objective** of each agent is to find the policy π_i^* , maximising the **expected discounted cumulative reward** considering the other agent's action once it reaches the game's Nash equilibrium

Agent Objectives at Nash Equilibrium

$$\pi_{\epsilon}^* = \arg \max_{\pi_{\epsilon}} \tilde{v}_i(s, \pi_{\epsilon}, \pi_{\alpha}^*)$$

$$\pi_{\alpha}^* = \arg \max_{\pi_{\alpha}} \tilde{v}_i(s, \pi_{\epsilon}^*, \pi_{\alpha})$$

Definition — Reward Machine (RM)

Finite state automaton designed to encode non-Markovian rewards with state transitions and output functions: $\mathcal{A} = (V, v_I, 2^{\mathcal{P}}, \mathbb{R}, \delta, \sigma)$

- **Input Alphabet ($2^{\mathcal{P}}$):** Interprets propositional logic variables $\mathcal{P}$ generated from lower-level environment updates.
- **Set of RM states (V):** Finite collection of internal tracking states, with v_I as the initial state.
- **Transition function:** Deterministic state progression $\delta : V \times 2^{\mathcal{P}} \rightarrow V$.
- **Output function:** Discrete scalars along active trajectory $\sigma : V \times 2^{\mathcal{P}} \rightarrow \mathbb{R}$.

Reward Machines example visualization

We define Reward Machines as a formalism to express **task specification**, decomposed into **several stages**.

Example — variant of the PAC-MAN game





Figure 1.,2. Variant of the PAC-MAN game is a symmetric zero-sum game where two agents try to capture each other. [1]

Connecting SG and RMs

A link between RM and SG is required

Labelling Function

$$L : S \times A_e \times A_a \times S \rightarrow 2^{\mathcal{P}}$$

$l_t = L(s_t, a_{a,t}, a_{e,t}, s_{t+1})$ is the set of **true propositions** in $\mathcal{P}$ at step t .

RM can encode non-Markovian RF of a stochastic game

Given a SG, the RM $\mathcal{A}_i$ **encodes** the non-markovian reward function R_i of the game, meaning that **for each trajectory** $s_0 a_{e,0} a_{a,0} s_1 \dots s_k a_{e,k} a_{a,k} s_{k+1}$ (i.e. $l_0 l_1 \dots, l_k$) it outputs the **reward sequence** $r_{i,0} r_{i,1} \dots, r_{i,k}$.

I.e. the RM maps **any sequence of events to a sequence of rewards**.

Q-learning with RM for SG

Problem formulation

The paper focuses on **two-agent general-sum stochastic games** (SG) with NM reward functions, where each agent is given a (different) task to complete.

Task completion **depends on the other's agent behavior**.

- Each agent as a separate reward machine.
- Agents operate in a shared and adversarial environment.

Then, the SG with non-markovian reward functions is formalized as **a product stochastic game** to obtain **Markovian rewards** using **reward machines**.

Markovian formulation: SGRM

Definition — Stochastic Game with Reward Machines (SGRM)

Given an SG $\mathcal{G}$ and RMs $\mathcal{A}_\epsilon, \mathcal{A}_\alpha$, a SGRM is defined as a product stochastic game:

$$\mathcal{H} = (S', s'_I, A_\epsilon, A_\alpha, p', R'_\epsilon, R'_\alpha, \gamma)$$

- **Augmented State Space:** $S' = S \times V_\epsilon \times V_\alpha$. An **augmented state** is $s' = (s, v_\epsilon, v_\alpha)$.
- **Transition Probabilities:** $p'(s', a_\epsilon, a_\alpha, \bar{s}') = p(s, a_\epsilon, a_\alpha, \bar{s})$ if $\bar{v}_i = \delta_i(v_i, L(s, a_\epsilon, a_\alpha, \bar{s}))$ for both $i \in \{\epsilon, \alpha\}$, and 0 otherwise.
- **Markovian Reward Functions:** $R'_i(s', a_\epsilon, a_\alpha, \bar{s}') = \sigma_i(v_i, L(s, a_\epsilon, a_\alpha, \bar{s}))$, $\forall i$

The **SGRM** is proved to have a Markovian **RF** to express the original **non-markovian RF** of the **SG**.

SGRM Steps

The agent consider S' to select actions: at **each time step simultaneous action selection and execution** occurs; this causes the **game state transition** from s to s' .

The SGRM labelling function L returns the **high-level event associated**, l_t . Using current label, the RM performs **RM state transition**, by moving from (v_e, v_a) to (v'_e, v'_a) , using each agent's specific RM, i.e. $v'_i = \delta(v_i, L(s, a_e, a_a, s'))$

The **transition of both game states and RM states** leads to the step **reward** for the agent $r_{i,t} = \sigma(v_i, L(s, a_e, a_a, s'))$

Nash Equilibrium and Q-functions in SGRM

Q-functions are used at Nash equilibrium.

QSG-RM

Q-function at a Nash equilibrium is the EDCR obtained by agent i when **both agents** follow a **joint** Nash Equilibrium strategy from the next period on.

So, it depends on **both agents actions**, and it's defined over the augmented state space.

Definition — Q-Function at a Nash Equilibrium in SGRM

Given agent i , given augmented state space $S \times V_e \times V_a$, we define

$$q_i^*(s, v_e, v_a, a_e, a_a) = r_i + \gamma \sum_{s', v'} p(s', v' | s, v, a_e, a_a) \tilde{v}_i(s', v'_e, v'_a, \pi_e^*, \pi_a^*)$$

Nash Equilibrium and Q-functions in SGRM (cont'd)

Using the q-function in the augmented space, best response policies are then originally extracted using the **Lemke-Howson method**.

By definition an agent strategy is optimal if the other agent behavior is considered: this requires both agents to **maintain two Q-functions**.

Agent learned (two) Q-functions

$q_{i,j}$ represents the Q-function for agent i as learned by agent j . It represents the **estimation that agent j does of agent i 's expected discounted cumulative reward**, assuming a game-theoretic baseline where each plays **accordingly to Nash Equilibrium** (i.e. choosing the best response to the other player -estimated-strategy.)

As a consequence, **each agent can estimate** (its own and) **other player's strategy** $\pi_{i,j}$, using the q-function estimation of its EDCR.

The Stage Game Matrix

Definition: Stage Game

A two-player stage game is defined as (r_ϵ, r_α) , where r_i is agent i 's reward function R_i , over the space of joint actions:

$$r_i = \{R_i(a_\alpha, a_\epsilon) | a_\alpha \in A_\alpha, a_\epsilon \in A_\epsilon\}$$

- In QRM-SG, a stage game is formulated at each time step based on current Q-functions in augmented state space.
- The **Lemke-Howson method** is used to derive best-response strategies at a Nash equilibrium.

QSG-RM Algorithm Overview

QRM-SG Algorithm Key aspects

Actor Mechanics (Action Selection)

- ϵ -greedy policy is adopted to balance the **exploration** and **exploitation**.

$$a_i \leftarrow \begin{cases} \operatorname{argmax}_{a \in A_i} \pi_{ji}(a \mid s, v_e, v_a) & \text{with probability } 1 - \epsilon \\ \text{random} & \text{with probability } \epsilon \end{cases}$$

Critic Updates (Learning Loop)

Values are updated according to standard stochastic approximation:

$$q_{ji} \leftarrow (1 - \alpha_t) q_{ji} + \alpha_t (r_{i,t} + \gamma \bar{q}_{ji}(s', v'_e, v'_a))$$

where

$$\bar{q}_{ji} \leftarrow \pi_{e,i}(s', v'_e, v'_a) \pi_{a,i}(s', v'_e, v'_a) q_{ji}(s', v'_e, v'_a) \in \mathbb{R}$$

QRM-SG Phase 1 — Initialization

Process Breakdown (Lines 1–4)

Hyperparameter: episode length $eplength$, γ , ϵ

Input: Reward machines $\mathcal{A}_\epsilon, \mathcal{A}_\alpha$

$s \leftarrow InitialState(); v_\epsilon \leftarrow v_{I,\epsilon}; v_\alpha \leftarrow v_{I,\alpha}$

$q_{\epsilon i}(\cdot), q_{\alpha i}(\cdot) \leftarrow InitialQfunctions()$

Analysis:

- Each agent initializes **two** Q-tables to store values for self and opponent, allowing local equilibrium computation.
- A SGRM (augmented) states are defined in the product space $S \times V_\epsilon \times V_\alpha$; each tuple element s , v_ϵ and v_α is **individually initialized**.

QRM-SG Phase 2 — Nash-Action Selection

Process Breakdown (Lines 7–13)

```
# for each episode, for each time step  $t$ :  
 $\bar{\epsilon} \leftarrow \text{GenerateRandomValue}(0, 1)$   
if  $\bar{\epsilon} < \epsilon$  then  
     $a_i \leftarrow \text{ChooseActionRandomly}()$  # [Up, Down, Left, Right]  
else  
    for  $i \in \{\epsilon, \alpha\}$  do  
         $\pi_{\epsilon i}, \pi_{\alpha i} \leftarrow \text{CalculateNashEquilibrium}(q_{\epsilon i}, q_{\alpha i})$   
         $a_i \leftarrow \operatorname{argmax} \pi_{ii}(a|s, v_{\epsilon}, v_{\alpha})$   
    end for  
end if
```

At **each step**, agents solve a bimatrix stage game using the **Lemke-Howson method** to find mixed strategy equilibria.

QRM-SG Phase 3 — Reward Machine Dynamics

Previously selected step action gets executed:

Process Breakdown (Lines 14–18)

$s' \leftarrow \text{ExecuteAction}(s, a_e, a_a) \#$ SG state transition, i.e. new coords

$l_t \leftarrow L(s, a_e, a_a, s')$

for $i \in \{e, a\}$ **do**

$v'_i \leftarrow \delta_i(v_i, l_t); r_{i,t} \leftarrow \sigma_i(v_i, l_t)$

end for

- L bridges low-level physical dynamics ($s \rightarrow s'$) with high-level automata state tracking, detecting high-level environmental events. (e.g., "Agent i did X").
- Rewards $r_{i,t}$ are generated by the RM transitions, making them **context/history aware**.

Nash Q-Learning

Given an agent in a one player game setting, **off policy Q-learning** method can be used to update the Q-function estimation, as a TD learning method, using the following update rule:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

But the max operator is **not applicable in a multi-agent setting**, where reward depends on joint behavior. This is then substituted by the **Nash equilibrium**, for each agent q-function:

$$Q_i(s, \mathbf{a}) \leftarrow Q_i(s, \mathbf{a}) + \alpha[r_i + \gamma \text{Nash}(Q_i(s', \mathbf{a})) - Q_i(s, \mathbf{a})]$$

Where $\text{Nash}(Q_i(s', \mathbf{a}))$ is the **expected value of agent i** when both agents follow the Nash strategy profile **at the next state** (s', v'_e, v'_a) :

$$\text{Nash}(Q_i(s', \mathbf{a})) = \mathbb{E}_{\mathbf{a}}[Q_i(s', \mathbf{a})] = \pi_{ei}(s', v'_e, v'_a) \pi_{ai}(s', v'_e, v'_a) Q_i(s', v'_e, v'_a)$$

Nash Q-Learning (cont'd) and Q updates

This means that, just like in standard Q-learning, the algorithm **looks ahead** to the next state (s', v'_e, v'_a) , and **instead of looking at the max entry** with respect to all actions, it looks at **all state-action pairs and considers the average value**, weighted by **relative action choice probabilities** at the Nash equilibrium.

Process Breakdown (Lines 19–24)

$$\pi_{ji}(\cdot | s', v') \leftarrow \text{Nash}(q_{ei}(s', v'), q_{ai}(s', v'))$$

$$\bar{q}_{ji} \leftarrow \pi_{ei} \pi_{ai} q_{ji}(s', v')$$

$$q_{ji} \leftarrow (1 - \alpha_t) q_{ji} + \alpha_t (r_{j,t} + \gamma \bar{q}_{ji})$$

Theoretical Foundations

Convergence Assumptions

Assumption 1: Exploration

Every state $s \in S$, RM states $v_e \in V_e$, $v_a \in V_a$, and actions $a_e \in A_e$, $a_a \in A_a$, are visited infinitely often when the number of episodes goes to infinity.

Assumption 2: Learning Rate

The learning rate α_t satisfies the following conditions for all t :

1. $0 < \alpha_t < 1$, $\sum_t \alpha_t(s, v_e, v_a, a_e, a_a) = \infty$, $\sum_t [\alpha_t(s, v_e, v_a, a_e, a_a)]^2 < \infty$, and the latter two hold uniformly and with probability 1.
2. $\alpha_t(s, v_e, v_a, a_e, a_a) = 0$ if $(s, v_e, v_a, a_e, a_a) \neq (s^t, v_e^t, v_a^t, a_e^t, a_a^t)$.

Assumption 3: Stage Game Structure

One of the following conditions holds during learning:

1. Every stage game $(q_{ei}(s, v_e, v_a), q_{ai}(s, v_e, v_a))$, $i \in \{e, a\}$, for all t, s, v_e , and v_a , has a **global optimal point** $(\tilde{\pi}_{ei}, \tilde{\pi}_{ai})$ (i.e., $\tilde{\pi}_{ji}q_{ji} \geq \tilde{\pi}'_{ji}q_{ji}$ for any $\tilde{\pi}'_{ji}$, $j \in \{e, a\}$) and agents' rewards in this equilibrium are used to update their Q-functions.
2. Every stage game $(q_{ei}(s, v_e, v_a), q_{ai}(s, v_e, v_a))$, $i \in \{e, a\}$, for all t, s, v_e , and v_a , has a **saddle point** $(\tilde{\pi}_{ei}, \tilde{\pi}_{ai})$ (i.e., $\tilde{\pi}_{ji}\tilde{\pi}_{-ji}q_{ji} \geq \tilde{\pi}'_{ji}\tilde{\pi}_{-ji}q_{ji}$ for any $\tilde{\pi}'_{ji}$, $\tilde{\pi}_{ji}\tilde{\pi}_{-ji}q_{ji} \leq \tilde{\pi}_{ji}\tilde{\pi}'_{-ji}q_{ji}$ for any $\tilde{\pi}'_{-ji}$, $j \in \{e, a\}$), and agents' rewards in this equilibrium are used to update their Q-functions.

Theorem 1: Convergence Guarantee

Theorem 1

Under Assumptions 1, 2 and 3, the sequence $q_{it} = (q_{\epsilon i}^t, q_{\alpha i}^t)$ at time t , $i \in \{\epsilon, \alpha\}$ updated by:

$$q_{ji}(s, v, a) = (1 - \alpha_t)q_{ji}(s, v, a) + \\ \alpha_t(r_{j,t} + \gamma \pi_{\epsilon i}(\cdot | s', v') \pi_{\alpha i}(\cdot | s', v') q_{ji}(s', v'))$$

for $j \in \{\epsilon, \alpha\}$, where $(\pi_{\epsilon i}(\cdot | s', v'), \pi_{\alpha i}(\cdot | s', v'))$ is the appropriate type of Nash equilibrium solution for the stage game $(q_{\epsilon i}^t(s', v'), q_{\alpha i}^t(s', v'))$, converges to the Q-functions at a Nash equilibrium $(q_{\epsilon i}^*, q_{\alpha i}^*)$.

Algorithm 1: QRM-SG Full Pseudocode

```
1: Initialize  $q_{ei}$  and  $q_{ai}$  for all  $i \in \{e, a\}$ 
2: for episode  $l = 1, 2, \dots$  do
3:    $s \leftarrow s_l, v_e \leftarrow v_{l,e}, v_a \leftarrow v_{l,a}$ 
4:   for  $t = 0$  to  $eplength$  do
5:     Choose action  $a_i$  using  $\epsilon$ -greedy policy and Lemke-Howson
6:     Execute actions, observe  $s'$ , and high-level event  $l_t = L(s, a_e, a_a, s')$ 
7:     Update RM states:  $v'_i \leftarrow \delta_i(v_i, l_t)$  and reward  $r_{i,t} \leftarrow \sigma_i(v_i, l_t)$ 
8:     Compute Nash strategy  $\pi_{ji}(\cdot | s', v')$  at next augmented state
9:     Update Q-functions:
10:     $q_{ji} \leftarrow (1 - \alpha_t)q_{ji} + \alpha_t(r_{j,t} + \gamma \sum \pi \pi q'_{ji})$ 
11:     $s \leftarrow s', v_e \leftarrow v'_e, v_a \leftarrow v'_a$ 
12:   end for
13: end for
```

Algorithm Analysis and Convergence Properties

- **Off-Policy Learning:** RMs enable decomposition of complex RL problems into structured subproblems for efficient learning.
- **Coupled Learning:** The learning processes of agents are coupled because the high-level events depend on joint actions.
- **Decentralized View:** Each agent maintains a "centralized" perspective by learning everyone's Q-functions to compute equilibria locally.

Mathematical Foundation: Convergence Proof

The proof relies on viewing the update as a stochastic approximation of a mapping F^t :

Key Lemmas from Supplementary Materials

- **Lemma 2:** Establishes general convergence for iterations $q^{t+1} = (1 - \alpha_t)q^t + \alpha_t[F^t q^t]$.
- **Lemma 4:** Shows that Q-functions at a Nash equilibrium (q^*) are fixed points of F^t , i.e., $q^* = \mathbb{E}[F^t q^*]$.
- **Lemma 7:** Proves F^t is a **contraction mapping** under Assumption 3:

$$\|F^t q - F^t \dot{q}\| \leq \gamma \|q - \dot{q}\|$$

Results and Applications

- **Testing Environment:** Evaluated on 6×6 (and 12×12) competitive, sparse-reward grid world variations of Pac-Man.
- **Comparative Baselines:**
 - Nash-Q: Pure spatial state-action tracking.
 - Nash-QAS: Incorporates an unstructured history status vector.
 - MADDPG-SG & MADDPG-AS: Deep multi-agent deterministic policy gradient variants.

Case Study Empirical Analysis

- **Case Study I (Sequential Ordering):** QRM-SG converges to equilibrium by $\approx 7,500$ episodes, whereas unaugmented baselines fail entirely.
- **Case Study II & III (Energy Loops):** Explores complex recharge loops with randomized origin points. QRM-SG reaches stability within $\approx 1,000$ to $1,500$ runs.
- **Scalability:** Successfully finds the Nash equilibrium on an expanded 12×12 map space by episode 21,000.

Conclusion

Summary & Future Outlook

- **Core Contribution:** Bridges reward machines with MARL to model non-Markovian constraints in stochastic games.
- **Robust Architecture:** Augmented state transformations allow tabular actors to navigate complex environments successfully.
- **Future Directions:** Dynamically learning unknown reward machine configurations simultaneously during policy calculation; extensions to general-sum games with more agents.

References

- [1] Jueming Hu, Jean-Raphaël Gaglione, Yanze Wang, Zhe Xu, Ufuk Topcu, and Yongming Liu.

Reinforcement learning with reward machines in stochastic games.

arXiv preprint arXiv:2305.17372v3, cs.MA, 2023.